: cenzth@gmail.com
: +44 7803453314

Cenz, Wong Tsz Ho

WORK EXPERIENCES

Data Engineer / Toyota Connected Europe / Dec 2024 – Now

- **AI-Native Engineering:** Standardised codebase architecture to be "LLM-readable," optimising for semantic clarity and context window efficiency; reduced token consumption by restructuring legacy patterns into modular, AI-ready components.
- **Agentic Automation:** Spearheading the internal transition toward autonomous operations by building prototypes for task-automation agents designed to handle repetitive engineering toil (e.g., automated PR analysis and documentation updates).
- **Semantic Data Quality:** Developed a custom framework for validating telematics data using semantic reasoning, moving beyond deterministic schema checks to ensure data integrity for downstream ML models.
- **Internal Tooling & DX:** Engineered a suite of internal developer tools to accelerate the telemetry pipeline lifecycle, prioritising modular design and auditability to ensure "production-grade" reliability across vehicle data streams.

Data Engineer / ASDA / Oct 2023 – Nov 2024

- **Graph-Augmented Analytics**: Spearheaded a Neo4j proof-of-concept to model complex customer relationships, evaluating graph-based approaches for downstream ML and identity resolution.
- **Engineering Excellence**: Developed and maintained production-grade internal Python packages to standardise ETL workflows, improving deployment velocity and code reusability across the data org.
- **Identity Resolution**: Engineered end-to-end Spark pipelines for customer data unification, implementing automated data quality frameworks and deterministic validation suites to ensure "gold-standard" record accuracy.
- **Infrastructure**: Established robust DevOps practices and Git-based CI/CD workflows, transitioning the team toward more rigorous software engineering standards.

Consultant in Data Engineering & Architecture / Ekimetrics / Dec 2021 – Aug 2023

- **LLM & RAG Prototyping:** Built early-stage Proof-of-Concepts using LangChain and Vector Databases, focusing on document retrieval and context injection for enterprise-specific use cases.
- **Global Scalability:** Architected and deployed multi-region data lakehouses on Azure Databricks for a global automotive lead, synchronising data across 17 countries while enforcing strict GDPR and governance protocols.
- **Serverless Orchestration**: Designed AWS-native ingestion engines using Lambda and containerised microservices to handle high-volume event streams with integrated automated quality checks.
- **MLOps**: Led the development of end-to-end ML pipelines on Databricks, processing multi-million record datasets for advanced segmentation and predictive modeling.

Engr Project Specialist (Innovation Lab) / Hewlett Packard Enterprise / Jun 2020 – Dec 2021

- **Enterprise Lakehouse Consulting**: Provided technical architecture guidance on HPE Ezmeral, demonstrating cloud-native data orchestration and hybrid-cloud ML capabilities.
- **Distributed Systems**: Developed data pipelines utilising the Hadoop ecosystem (Hive, Kafka, Spark) to support large-scale enterprise analytics.

EDUCATION

AI Performance Engineering / Nebius Academy / 2026 – Now

✓ Agents, LLM architecture, MLOps, Performance engineering, LLM post-training

MSc in Big Data Technology / The Hong Kong University of Science and Technology / 2020 – 2022

✓ Distributed Computing, Parallel Computing, Natural Language Processing, Machine Learning, Time Series Analysis, Data Mining

BEng(Hons) in Electronic and Information Engineering / The Hong Kong Polytechnic University / 2015 – 2020

✓ Machine Learning and IoT, Circuit Analysis and Logic Design, Embedded System Programming, Microcontroller System and Interface Design